

CS 344 Operating Systems Laboratory Assignment 1

TEAM MEMBERS:

NAME	ROLL NUMBER
RITWIK GANGULY	180101067
PULKIT CHANGOIWALA	180101093
SAMAY VARSHNEY	180101097
UDANDARAO SAI SANDEEP	180123063

EXERCISE 1:

Modified Program **ex1.c** which now includes inline assembly that increments the value of x by 1.

```
#include<stdio.h>
int main(int argc, char **argv)
{
    int x=1;
    printf("Hello x = %d\n", x);
    asm("incl %0": "+r"(x));
    /*
    First method to increment the value of x using inline assembly language having one line
    asm("incl %0": "+r"(x));
    1) incl command adds 1 to the 32-bit contents of the variable specified
    2) %0 refers to the first variable passed (which in this case is x)
    3) the '+' sign before r denotes that x acts as both input and output
    */
    /*
    Second method to increment the value of x using inline assembly language having
    multiple lines (uncomment to use this)
    asm("mov %1, %0 \n\t"
        "add $1, %0"
        : "=r" (x)
        : "r" (x));
    */
    printf("Hello x = %d after increment\n", x);
    if(x == 2) {
        printf("OK\n");
    }
    else {
        printf("ERROR\n");
    }
}
```

EXERCISE 2:

We are trying to explain the instructions to guess what BIOS might be trying to do:

1st instruction: [f000:fff0] 0xffff0: jmp \$0x3630,\$0xf000e05b

- Jump to CS = \$0xf000 & IP = 0xe05b
- 0x3630 is jump to this CS (earlier in the BIOS)
- 0xf000e05b is the IP which is different from the lab because it is 32 bits rather than 16 bits and that is all the way into the top of the extended memory location but before the memory mapped PCI device location reserved by the BIOS

2nd Instruction: [f000:e05b] 0xfe05b: cmpw \$0xffc8,%cs:(%esi)

- Compare content at 0xffc8 & with content at code segment offset with value at esi.
- esi:- 32-bit source index register

3rd Instruction: [f000:e062] 0xfe062: jne 0xd241d0b2

- Jump to 0xd241d0b2 if the above comparison does not set ZF

4th instruction: [f000:e066] 0xfe066: xor %edx,%edx

- ZF was set thus jump of previous instruction doesn't occur
- It set edx to zero, edx is 32-bit general-purpose register.

5th instruction: [f000:e068] 0xfe068: mov %edx,%ss

- Move content of stack segment register(ss) to edx

6th instruction: [f000:e06a] 0xfe06a: mov \$0x7000,%sp

- Move content at the location pointed 16-bit stack pointer(sp) to \$0x7000

```
The target architecture is assumed to be i8086
[f000:fff0] 0xffff0: jmp $0x3630,$0xf000e05b
0x0000fff0 in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration
of GDB. Attempting to continue with the default i8086 settings.

(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) si
[f000:e062] 0xfe062: jne 0xd241d416
0x0000e062 in ?? ()
(gdb) si
[f000:e066] 0xfe066: xor %edx,%edx
0x0000e066 in ?? ()
(gdb) si
[f000:e068] 0xfe068: mov %edx,%ss
0x0000e068 in ?? ()
(gdb) si
[f000:e06a] 0xfe06a: mov $0x7000,%sp
0x0000e06a in ?? ()
```

EXERCISE 3:

Comparison of code at 0x7C00 memory location using first few instructions between original boot loader source code (bootasm.S) and GDB and bootblock.asm :

In Bootasm.S:

```
.code16                                # Assemble for 16-bit mode
.globl start
start:
    cli                                # BIOS enabled interrupts; disable
    xorw %ax,%ax                       # Set %ax to zero
    movw %ax,%ds                       # -> Data Segment
    movw %ax,%es                       # -> Extra Segment
```

```

        movw %ax,%ss          # -> Stack Segment
seta20.1:
        inb  $0x64,%al        # Wait for not busy
        testb $0x2,%al
        jnz  seta20.1
        movb $0xd1,%al        # 0xd1 -> port 0x64
        outb %al,$0x64

```

In Bootblock.asm

```

.code16                          # Assemble for 16-bit mode
.globl start
start:
    cli                          # BIOS enabled interrupts; disable
        7c00: fa                cli
    # Zero data segment registers DS, ES, and SS.
    xorw %ax,%ax                # Set %ax to zero
        7c01: 31 c0            xor  %eax,%eax
    movw %ax,%ds                # -> Data Segment
        7c03: 8e d8            mov  %eax,%ds
    movw %ax,%es                # -> Extra Segment
        7c05: 8e c0            mov  %eax,%es
    movw %ax,%ss                # -> Stack Segment
        7c07: 8e d0            mov  %eax,%ss

00007c09 <seta20.1>:
seta20.1:
    inb  $0x64,%al              # Wait for not busy
        7c09: e4 64            in   $0x64,%al
    testb $0x2,%al
        7c0b: a8 02            test $0x2,%al
    jnz  seta20.1
        7c0d: 75 fa            jne  7c09 <seta20.1>

    movb $0xd1,%al              # 0xd1 -> port 0x64
        7c0f: b0 d1            mov  $0xd1,%al
    outb %al,$0x64
        7c11: e6 64            out  %al,$0x64

```

In GDB:

```

(gdb) x/10i 0x7C00
0x7c00:    cli
=> 0x7c01:    xor    %eax,%eax
0x7c03:    mov    %eax,%ds
0x7c05:    mov    %eax,%es
0x7c07:    mov    %eax,%ss
0x7c09:    in     $0x64,%al
0x7c0b:    test   $0x2,%al
0x7c0d:    jne    0x7c09
0x7c0f:    mov    $0xd1,%al
0x7c11:    out    %al,$0x64

```

Tracing into Bootmain and Readsect:

Statements in readsect in bootmain.c:

- 1) waitdisk(); // Issue command
Assembly Instruction:
7c98: e8 e1 ff ff ff call 7c7e <waitdisk>
- 2) outb(0x1F2, 1); // count = 1
- 3) outb(0x1F3, offset);
- 4) outb(0x1F4, offset >> 8);
Assembly Instruction:
7cb0: 89 d8 mov %ebx,%eax
7cb2: c1 e8 08 shr \$0x8,%eax
7cb5: ba f4 01 00 00 mov \$0x1f4,%edx
7cba: ee out %al,(%dx)
- 5) outb(0x1F5, offset >> 16);
Assembly Instruction:
7cbb: 89 d8 mov %ebx,%eax
7cbd: c1 e8 10 shr \$0x10,%eax
7cc0: ba f5 01 00 00 mov \$0x1f5,%edx
7cc5: ee out %al,(%dx)
- 6) outb(0x1F6, (offset >> 24) | 0xE0);
Assembly Instruction:
7cc6: 89 d8 mov %ebx,%eax
7cc8: c1 e8 18 shr \$0x18,%eax
7ccb: 83 c8 e0 or \$0xffffffe0,%eax
7cce: ba f6 01 00 00 mov \$0x1f6,%edx
7cd3: ee out %al,(%dx)
7cd4: b8 20 00 00 00 mov \$0x20,%eax
7cd9: ba f7 01 00 00 mov \$0x1f7,%edx
7cde: ee out %al,(%dx)
- 7) outb(0x1F7, 0x20); // cmd 0x20 - read sectors
- 8) waitdisk(); // Read data.
Assembly Instruction:
7cdf: e8 9a ff ff ff call 7c7e <waitdisk>
- 9) insl(0x1F0, dst, SECTSIZE/4);

Begin and end of loop which reads remaining sectors:

```
for(; pa < epa; pa += SECTSIZE, offset++)  
    readsect(pa, offset);
```

What code will run after running out of loop:

```
readseg(pa, ph->filesz, ph->off);  
7d99: ff 73 04                pushl 0x4(%ebx)  
7d9c: ff 73 10                pushl 0x10(%ebx)  
7d9f: 57                    push %edi  
7da0: e8 53 ff ff ff        call 7cf8 <readseg>
```

Setting up breakpoint at 0x7d99 as it is the first instruction after running out of readseg function and loop consecutively.

```

(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) b *0x7d99
Breakpoint 2 at 0x7d99
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d99:      pushl   0x4(%ebx)

Thread 1 hit Breakpoint 2, 0x00007d99 in ?? ()
(gdb) x/10i 0x7d99
=> 0x7d99:      pushl   0x4(%ebx)
0x7d9c:      pushl   0x10(%ebx)
0x7d9f:      push    %edi
0x7da0:      call    0x7cf8
0x7da5:      mov     0x14(%ebx),%ecx
0x7da8:      mov     0x10(%ebx),%eax
0x7dab:      add     $0xc,%esp
0x7dae:      cmp     %eax,%ecx
0x7db0:      jbe     0x7d8f
0x7db2:      add     %eax,%edi

```

Answer the following questions:

a) The command `ljmp $(SEG_KCODE<<3), $start32` causes the switch from 16 to 32-bit mode in the bootasm.S which occurs at address **0x7C31**.

```

(gdb) b *0x7c29
Breakpoint 1 at 0x7c29
(gdb) c
Continuing.
[ 0:7c29] => 0x7c29:  mov     %eax,%cr0

Thread 1 hit Breakpoint 1, 0x00007c29 in ?? ()
(gdb) si
[ 0:7c2c] => 0x7c2c:  ljmp    $0xb866,$0x87c31
0x00007c2c in ?? ()
(gdb) si
The target architecture is assumed to be i386
=> 0x7c31:      mov     $0x10,%ax
0x00007c31 in ?? ()
(gdb) x/20i 0x7c29
0x7c29:      mov     %eax,%cr0
0x7c2c:      ljmp    $0xb866,$0x87c31
0x7c33:      adc     %al,(%eax)
0x7c35:      mov     %eax,%ds
0x7c37:      mov     %eax,%es
0x7c39:      mov     %eax,%ss
0x7c3b:      mov     $0x0,%ax
0x7c3f:      mov     %eax,%fs
0x7c41:      mov     %eax,%gs
0x7c43:      mov     $0x7c00,%esp
0x7c48:      call    0x7d3b
0x7c4d:      mov     $0x8a00,%ax
0x7c51:      mov     %ax,%dx
0x7c54:      out     %ax,(%dx)
0x7c56:      mov     $0x8ae0,%ax
0x7c5a:      out     %ax,(%dx)
0x7c5c:      jmp     0x7c5c
0x7c5e:      xchg    %ax,%ax
0x7c60:      add     %al,(%eax)
0x7c62:      add     %al,(%eax)

```

b) Last Instruction of boot loader executed:
in bootmain.c it is:
entry = (void*)(void))(elf->entry);
entry();

in bootblock.asm it is:
7d87: ff 15 18 00 01 00 call *0x10018

The First Instruction of Kernel it just loaded is:

0x10000c : mov %cr4,%eax

Also the first instruction of the kernel should be at **0x10018**.


```
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) b *0x7d87
Breakpoint 2 at 0x7d87
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d87:              call      *0x10018

Thread 1 hit Breakpoint 2, 0x00007d87 in ?? ()
(gdb) si
=> 0x10000c:          mov        %cr4,%eax
0x0010000c in ?? ()
(gdb) x/1x 0x10018
0x10018:              0x0010000c
```

c) The boot loader reads the number the program headers in the ELF header and loads them all. It finds this information in the ELF header.

```
ph = (struct proghdr*)((uchar*)elf + elf->phoff);
eph = ph + elf->phnum;
for(; ph < eph; ph++){
    pa = (uchar*)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
    if(ph->memsz > ph->filesz)
        stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
}
```



EXERCISE 4:

\$ objdump -h kernel

```
kernel:          file format elf32-i386
```

Sections:					
Idx	Name	Size	VMA	LMA	File off
0	.text	00006f12	80100000	00100000	00001000
		CONTENTS,	ALLOC, LOAD, READONLY, CODE		
1	.rodata	00000b6c	80106f20	00106f20	00007f20
		CONTENTS,	ALLOC, LOAD, READONLY, DATA		
2	.data	00002516	80108000	00108000	00009000
		CONTENTS,	ALLOC, LOAD, DATA		
3	.bss	0000af88	8010a520	0010a520	0000b516
		ALLOC			
4	.debug_line	000025f5	00000000	00000000	0000b516
		CONTENTS,	READONLY, DEBUGGING		
5	.debug_info	000105a6	00000000	00000000	0000db0b
		CONTENTS,	READONLY, DEBUGGING		
6	.debug_abbrev	0000397c	00000000	00000000	0001e0b1
		CONTENTS,	READONLY, DEBUGGING		
7	.debug_aranges	000003a8	00000000	00000000	00021a30
		CONTENTS,	READONLY, DEBUGGING		
8	.debug_str	00000e6f	00000000	00000000	00021dd8
		CONTENTS,	READONLY, DEBUGGING		
9	.debug_loc	00005294	00000000	00000000	00022c47
		CONTENTS,	READONLY, DEBUGGING		
10	.debug_ranges	00000700	00000000	00000000	00027edb
		CONTENTS,	READONLY, DEBUGGING		
11	.comment	00000029	00000000	00000000	000285db
		CONTENTS,	READONLY		

\$ objdump -h bootblock.o

```
bootblock.o:      file format elf32-i386
```

Sections:					
Idx	Name	Size	VMA	LMA	File off
0	.text	000001c0	00007c00	00007c00	00000074
		CONTENTS,	ALLOC, LOAD, CODE		
1	.eh_frame	000000bc	00007dc0	00007dc0	00000234
		CONTENTS,	ALLOC, LOAD, READONLY, DATA		
2	.comment	00000029	00000000	00000000	000002f0
		CONTENTS,	READONLY		
3	.debug_aranges	00000040	00000000	00000000	00000320
		CONTENTS,	READONLY, DEBUGGING		
4	.debug_info	0000050b	00000000	00000000	00000360
		CONTENTS,	READONLY, DEBUGGING		
5	.debug_abbrev	000001e3	00000000	00000000	0000086b
		CONTENTS,	READONLY, DEBUGGING		
6	.debug_line	0000012c	00000000	00000000	00000a4e
		CONTENTS,	READONLY, DEBUGGING		
7	.debug_str	000001d9	00000000	00000000	00000b7a
		CONTENTS,	READONLY, DEBUGGING		
8	.debug_loc	0000022a	00000000	00000000	00000d53
		CONTENTS,	READONLY, DEBUGGING		

Fields Explanation:

- 1) Name: Program Sections Name(Program Headers)
- 2) Size: Size of the loaded section
- 3) VMA: Link Address, The link address of a section is the memory address from which the section expects to execute.
- 4) LMA: Load Address, The load address of a section is the memory address at which that section should be loaded into memory.
- 5) File off: is this section's offset from the beginning of the file
- 6) Algn: It represents alignment
- 7) CONTENTS, ALLOC, LOAD, READONLY, DATA are flags. They represent that a particular section is to be LOADED or is READ ONLY.

EXERCISE 5:

When boot loader's link address is 0x7C00 then commands are running properly and transition from 16 to 32 bit was occurring at **0x7C31** address location as seen below:

```
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) b *0x7c31
Breakpoint 2 at 0x7c31
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7c31:      mov     $0x10,%ax

Thread 1 hit Breakpoint 2, 0x00007c31 in ?? ()
```

But when the boot loader's link address is changed to any other address (we took **0x7C24** in this case), after running

make clean

make

and restarting gdb

and continuing by putting breakpoint from address location 0x7C00,

then the boot loader is restarting again and again after running some instructions in the gdb.

```
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) b *0x7c55
Breakpoint 2 at 0x7c55
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) si
[ 0:7c01] => 0x7c01: xor     %eax,%eax
0x00007c01 in ?? ()
(gdb) si
[ 0:7c03] => 0x7c03: mov     %eax,%ds
0x00007c03 in ?? ()
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
```

As seen in the image above, we tried to run commands after continuing from breakpoint at 0x7C00 address location and we always end up hitting the same breakpoint at 0x7C00.

Also 16 to 32 bit architecture change didn't occurred as breakpoint **b *0x7C31** is not hitted which should be responsible for architecture change in this case.

ljmp \$(SEG_KCODE<<3), \$start32 is the first instruction that breaks.

Before changing the link address of the boot loader, from address 0x7C00, after performing 2-3 **si 10** instructions, architecture changed from 16 to 32 bit.

But after changing the link address to 0x7C24, architecture didn't change which means that the boot loader is not loaded properly at the changed link address.

\$ objdump -f kernel

```
kernel:      file format elf32-i386
architecture: i386, flags 0x00000112:
EXEC_P, HAS_SYMS, D_PAGED
start address 0x0010000c
```

EXERCISE 6:

At the point when BIOS enters the boot loader (at first breakpoint):

```
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) x/8i 0x00100000
0x100000: add    %al,(%eax)
0x100002: add    %al,(%eax)
0x100004: add    %al,(%eax)
0x100006: add    %al,(%eax)
0x100008: add    %al,(%eax)
0x10000a: add    %al,(%eax)
0x10000c: add    %al,(%eax)
0x10000e: add    %al,(%eax)
```

At the point when the boot loader enters the kernel (at second breakpoint):

```
(gdb) b *0x7d87
Breakpoint 2 at 0x7d87
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d87: call    *0x10018

Thread 1 hit Breakpoint 2, 0x00007d87 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x83e0200f
0x100010: 0x220f10c8 0x9000b8e0 0x220f0010 0xc0200fd8
(gdb) x/8i 0x00100000
0x100000: add    0x1bad(%eax),%dh
0x100006: add    %al,(%eax)
0x100008: decb   0x52(%edi)
0x10000b: in     $0xf,%al
0x10000d: and    %ah,%al
0x10000f: or     $0x10,%eax
0x100012: mov    %eax,%cr4
0x100015: mov    $0x109000,%eax
```

8 words of instruction at 0x00100000 at the point when BIOS enters the boot loader and 8 words of instruction at 0x00100000 at the point when the boot loader enters the kernel are different as when the BIOS enters and loads the boot loader, then it just loads it in memory location between 0x7C00 and 0x7DFF due to which all the 8 words of instructions are zero at 0x00100000. But when the boot loader enters the kernel, it already has performed the 16 to 32 bit transition and setting up of stack and also the boot loader loads kernel at memory locations including 0x00100000 which leads to new instructions at address 0x00100000.

EXERCISE 7:

In order to define our system call in xv6, we changed 5 files mentioned below.

1) **syscall.h:**

We added a new system call **#define SYS_wolfie 22** at 22nd position as 21 positions we already occupied by the inbuilt system calls in syscall.h.

2) **syscall.c:**

We added a pointer **[SYS_wolfie] sys_wolfie** to system call at 22nd position in syscall.c file in order to add our custom system call.

Then a function prototype **extern int sys_wolfie(void);** is added in syscall.c file which will be called by system call number 22.

3) **sysproc.c:**

System call function is implemented in sysproc.c.

```
int
sys_wolfie(void){
    // code is given in file sysproc.c
}
```

4) **usys.S:**

For creating an interface for your user program to access system call we added the following line in usys.S.

SYSCALL(wolfie)

5) **user.h:**

We added the following function that the user program will be calling in user.h.

int wolfie(void *buf, uint size);

Call to the above function from the user program will be simply mapped to system call number 22 which is defined as **SYS_wolfie** preprocessor directive. The system knows what exactly is this system call and how to handle it.

EXERCISE 8:

We save a C program named **wolfietest.c** inside the source code directory of xv6 operating system.

Then we edit the MakeFile and added below changes in MakeFile:

ls command listing all files present inside fs.img

```
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200
init: starting sh
$ ls
.          1 1 512
..         1 1 512
README    2 2 2286
cat        2 3 13640
echo       2 4 12644
forktest   2 5 8084
grep       2 6 15516
init       2 7 13232
kill       2 8 12700
ln         2 9 12600
ls         2 10 14784
mkdir      2 11 12784
rm         2 12 12760
sh         2 13 23248
stressfs   2 14 13428
usertests  2 15 56360
wc         2 16 14180
zombie     2 17 12424
wolfietest 2 18 12740
console    3 19 0
```